

# PRÁCTICA MONGODB

Hacer una prueba de rendimiento con y sin índices, justificando el resultado de:

1. Operaciones de búsqueda.
2. Operaciones de agregación.
3. Operaciones de actualización.
4. Operaciones de borrado.

Descripción de qué es un índice, cuándo utilizarlo, etc.

Selección de un dataset en el repositorio de GitHub y carga en MongoDB.

```
Imf > db.practicafinal.insertMany([{"_id":0,"name":"aimee
Zank","scores":[{"score":1.463179736705023,"type":"exam"},{"score":11.7827330995
7772,"type":"quiz"},{"score":35.8740349954354,"type":"homework"}]},
... {"_id":1,"name":"Aurelia
Menendez","scores":[{"score":60.06045071030959,"type":"exam"},{"score":52.797906
91903873,"type":"quiz"},{"score":71.76133439165544,"type":"homework"}]},
... {"_id":2,"name":"Corliss
Zuk","scores":[{"score":67.03077096065002,"type":"exam"},{"score":6.301851677835
235,"type":"quiz"},{"score":66.28344683278382,"type":"homework"}]},
... {"_id":3,"name":"Bao
Ziglar","scores":[{"score":71.64343899778332,"type":"exam"},{"score":24.8022129365
0313,"type":"quiz"},{"score":42.26147058804812,"type":"homework"}]},
... {"_id":4,"name":"Zachary
Langlais","scores":[{"score":78.68385091304332,"type":"exam"},{"score":90.29631013
68042,"type":"quiz"},{"score":34.41620148042529,"type":"homework"}]},
... {"_id":5,"name":"Wilburn
Spiess","scores":[{"score":44.87186330181261,"type":"exam"},{"score":25.723951146
68016,"type":"quiz"},{"score":63.42288310628662,"type":"homework"}]},
... {"_id":6,"name":"Jenette
Flanders","scores":[{"score":37.32285459166097,"type":"exam"},{"score":28.3263497
6913737,"type":"quiz"},{"score":81.57115318686338,"type":"homework"}]},
... {"_id":7,"name":"Salena
Olmos","scores":[{"score":90.37826509157176,"type":"exam"},{"score":42.487806669
56811,"type":"quiz"},{"score":96.52986171633331,"type":"homework"}]},
... {"_id":8,"name":"Daphne
Zheng","scores":[{"score":22.13583712862635,"type":"exam"},{"score":14.639699413
35069,"type":"quiz"},{"score":75.94123677556644,"type":"homework"}]}],
```

```
... {"_id":9,"name":"Sanda
Ryba","scores":[{"score":97.00509953654694,"type":"exam"},{"score":97.8044963253
8915,"type":"quiz"},{"score":25.27368532432955,"type":"homework"}]},

... {"_id":10,"name":"Denisha
Cast","scores":[{"score":45.61876862259409,"type":"exam"},{"score":98.35723209418
343,"type":"quiz"},{"score":55.90835657173456,"type":"homework"}]},

... {"_id":11,"name":"Marcus
Blohm","scores":[{"score":78.42617835651868,"type":"exam"},{"score":82.583728179
30675,"type":"quiz"},{"score":87.49924733328717,"type":"homework"}]},

... {"_id":12,"name":"Quincy Danaher"
,"scores":[{"score":54.29841278520669,"type":"exam"},{"score":85.61270164694737,"
type":"quiz"},{"score":80.40732356118075,"type":"homework"}]},

... {"_id":13,"name":"Jessika
Dagenais","scores":[{"score":90.47179954427436,"type":"exam"},{"score":90.3001402
468489,"type":"quiz"},{"score":95.17753772405909,"type":"homework"}]},

... {"_id":14,"name":"Alix
Sherrill","scores":[{"score":25.15924151998215,"type":"exam"},{"score":68.644840476
92098,"type":"quiz"},{"score":24.68462152686763,"type":"homework"}]},

... {"_id":15,"name":"Tambra
Mercure","scores":[{"score":69.1565022533158,"type":"exam"},{"score":3.311794422
000724,"type":"quiz"},{"score":45.03178973642521,"type":"homework"}]},

... {"_id":16,"name":"Dodie
Staller","scores":[{"score":7.772386442858281,"type":"exam"},{"score":31.843002351
04542,"type":"quiz"},{"score":80.52136407989194,"type":"homework"}]},

... {"_id":17,"name":"Fletcher
Mcconnell","scores":[{"score":39.41011069729274,"type":"exam"},{"score":81.132703
07809924,"type":"quiz"},{"score":97.70116640402922,"type":"homework"}]},

... {"_id":18,"name":"Verdell
Sowski","scores":[{"score":62.12870233109035,"type":"exam"},{"score":84.7458622
0889356,"type":"quiz"},{"score":81.58947824932574,"type":"homework"}]},

... {"_id":19,"name":"Gisela
Levin","scores":[{"score":44.51211101958831,"type":"exam"},{"score":0.65784979663
68002,"type":"quiz"},{"score":93.36341655949683,"type":"homework"}]},

... {"_id":20,"name":"Tressa
Schwing","scores":[{"score":42.17439799514388,"type":"exam"},{"score":71.99314840
599558,"type":"quiz"},{"score":81.23972632069464,"type":"homework"}]},
```

```
... {"_id":21,"name":"Rosana
Vales","scores":[{"score":46.2289476258328,"type":"exam"}, {"score":98.34164225207
036,"type":"quiz"}, {"score":36.18769746805938,"type":"homework"}]},

... {"_id":22,"name":"Margart
Vitello","scores":[{"score":75.04996547553947,"type":"exam"}, {"score":10.230464758
99236,"type":"quiz"}, {"score":96.72520512117761,"type":"homework"}]},

... {"_id":23,"name":"Tamika
Schildgen","scores":[{"score":45.65432764125526,"type":"exam"}, {"score":64.3292704
9658846,"type":"quiz"}, {"score":83.53933351660562,"type":"homework"}]},

... {"_id":24,"name":"Jesusa
Rickenbacker","scores":[{"score":86.0319702155683,"type":"exam"}, {"score":1.967495
200433389,"type":"quiz"}, {"score":61.10861071547914,"type":"homework"}]}

Así hasta _id=199
```

## **1. ANÁLISIS DE LA ESTRUCTURA DEL DATASET SELECCIONADO**

MongoDB almacena los documentos en un formato llamado BSON (o JSON binario), siendo cada documento descriptivo, en el que incluye su longitud y los tipos de datos.

La estructura de datos está compuesta por colecciones y documentos, como una sucesión de campos en forma de clave:valor

MongoDB ha almacenado los documentos en colecciones. Cada documento almacenado en una colección, tiene un campo `_id` único que identifica inmutablemente a dicho documento, actuando como clave primaria (ej.: `_id:0`, `_id:1` .... `_id:199`), y a este campo le asignamos el valor (ej.: `"name":"Aimee Zank"`, `"name":"Aurelia Menendez"`...`"name":"Rae Kohout"`).

El campo también tiene un subconjunto de objetos, como es el caso del campo `"scores"`, que es a su vez un array, formado por una serie de diccionarios.

```

    ]
  }
  > db.practicafinal.find({name:"Nobuko Linzey"}).pretty()
{
  "_id" : 110,
  "name" : "Nobuko Linzey",
  "scores" : [
    {
      "score" : 67.40792606687442,
      "type" : "exam"
    },
    {
      "score" : 58.58331128403415,
      "type" : "quiz"
    },
    {
      "score" : 47.44831568815929,
      "type" : "homework"
    }
  ]
}
>

```

## 2. PRUEBA DE RENDIMIENTO

Primero veremos el rendimiento sin índices:

### 3.1. Operaciones de búsqueda.

#### SIN INDICE

```

imf> db.practicafinal.find({name:"Nobuko Linzey"},{ id:0})
[
  {
    name: 'Nobuko Linzey',
    scores: [
      { score: 80.5, type: 'exam' },
      { score: 58.58331128403415, type: 'quiz' },
      { score: 47.44831568815929, type: 'homework' }
    ]
  }
]

```

Vemos ahora el rendimiento con `explain("executionStats")`

```

lmf> db.practicafinal.find({name:"Nobuko Linzey"},{_id:0}).explain("executionStats")
{
  explainVersion: '1',
  queryPlanner: {
    namespace: 'lmf.practicafinal',
    indexFilterSet: false,
    parsedQuery: { name: { '$eq': 'Nobuko Linzey' } },
    queryHash: '64908032',
    planCacheKey: '64908032',
    maxIndexedOrSolutionsReached: false,
    maxIndexedAndSolutionsReached: false,
    maxScansToExplodeReached: false,
    winningPlan: {
      stage: 'PROJECTION_DEFAULT',
      transformBy: { _id: 0 },
      inputStage: {
        stage: 'COLLSCAN',
        filter: { name: { '$eq': 'Nobuko Linzey' } },
        direction: 'forward'
      }
    },
    rejectedPlans: []
  },
  executionStats: {
    executionSuccess: true,
    nReturned: 1,
    executionTimeMillis: 0,
    totalKeysExamined: 0,
    totalDocsExamined: 200,
    executionStages: {
      stage: 'PROJECTION_DEFAULT',
      nReturned: 1,
      executionTimeMillisEstimate: 0,
      works: 202,
      advanced: 1,
      needTime: 200,
      needYield: 0,
      saveState: 0,
      restoreState: 0,
      isEOF: 1,
      transformBy: { _id: 0 },
      inputStage: {
        stage: 'COLLSCAN',
        filter: { name: { '$eq': 'Nobuko Linzey' } },
        nReturned: 1,
        executionTimeMillisEstimate: 0,
        works: 202,

```

```

        advanced: 1,
        needTime: 200,
        needYield: 0,
        saveState: 0,
        restoreState: 0,
        isEOF: 1,
        direction: 'forward',
        docsExamined: 200
      }
    },
    command: {
      find: 'practicafinal',
      filter: { name: 'Nobuko Linzey' },
      projection: { _id: 0 },
      '$db': 'lmf'
    },
    serverInfo: {
      host: '13e6561c38af',
      port: 27017,
      version: '6.0.1',
      gitVersion: '32f0f9c88dc44a2c8073a5bd47cf779d4bfdee6b'
    },
    serverParameters: {
      internalQueryFacetBufferSizeBytes: 104857600,
      internalQueryFacetMaxOutputDocSizeBytes: 104857600,
      internalLookupStageIntermediateDocumentMaxSizeBytes: 104857600,
      internalDocumentSourceGroupMaxMemoryBytes: 104857600,
      internalQueryMaxBlockingSortMemoryUsageBytes: 104857600,
      internalQueryProhibitBlockingMergeOnMongoS: 0,
      internalQueryMaxAddToSetBytes: 104857600,
      internalDocumentSourceSetWindowFieldsMaxMemoryBytes: 104857600
    },
    ok: 1
  }
}

```

Lo importante aquí:

nReturned: 1 (devuelve una salida)

TotalKeysExamined: 0 (no usa índice)

totalDocsExamined: 200 (examinó 200 documentos y devolvió 1)

## CON INDICE

```
imf> db.practicafinal.createIndex( { name: 1 } )
```

name\_1

🔗 Hemos creado índice en **name**

Vemos rendimiento:

```
imf> db.practicafinal.find({name: "Nobuko Linzey"},{_id:0}).explain("executionStats")
{
  explainVersion: '1',
  queryPlanner: {
    namespace: 'imf.practicafinal',
    indexFilterSet: false,
    parsedQuery: { name: { '$eq': 'Nobuko Linzey' } },
    queryHash: '64908032',
    planCacheKey: 'A6C0273F',
    maxIndexedOrSolutionsReached: false,
    maxIndexedAndSolutionsReached: false,
    maxScansToExplodeReached: false,
    winningPlan: {
      stage: 'PROJECTION_DEFAULT',
      transformBy: { _id: 0 },
      inputStage: {
        stage: 'FETCH',
        inputStage: {
          stage: 'IXSCAN',
          keyPattern: { name: 1 },
          indexName: 'name_1',
          isMultiKey: false,
          multiKeyPaths: { name: [] },
          isUnique: false,
          isSparse: false,
          isPartial: false,
          indexVersion: 2,
          direction: 'forward',
          indexBounds: { name: [ '['Nobuko Linzey', 'Nobuko Linzey']' ] }
        }
      }
    },
    rejectedPlans: []
  },
  executionStats: {
    executionSuccess: true,
    nReturned: 1,
    executionTimeMillis: 2,
    totalKeysExamined: 1,
    totalDocsExamined: 1,
    executionStages: {
      stage: 'PROJECTION_DEFAULT',
      nReturned: 1,
      executionTimeMillisEstimate: 0,
      works: 2,
      advanced: 1,
      needTime: 0,
      needYield: 0,

```

```

    saveState: 0,
    restoreState: 0,
    isEOF: 1,
    transformBy: { _id: 0 },
    inputStage: {
      stage: 'FETCH',
      nReturned: 1,
      executionTimeMillisEstimate: 0,
      works: 2,
      advanced: 1,
      needTime: 0,
      needYield: 0,
      saveState: 0,
      restoreState: 0,
      isEOF: 1,
      docsExamined: 1,
      alreadyHasObj: 0,
      inputStage: {
        stage: 'IXSCAN',
        nReturned: 1,
        executionTimeMillisEstimate: 0,
        works: 2,
        advanced: 1,
        needTime: 0,
        needYield: 0,
        saveState: 0,
        restoreState: 0,
        isEOF: 1,
        keyPattern: { name: 1 },
        indexName: 'name_1',
        isMultiKey: false,
        multiKeyPaths: { name: [] },
        isUnique: false,
        isSparse: false,
        isPartial: false,
        indexVersion: 2,
        direction: 'forward',
        indexBounds: { name: [ ["Nobuko Linzey", "Nobuko Linzey"] ] },
        keysExamined: 1,
        seeks: 1,
        dupsTested: 0,
        dupsDropped: 0
      }
    }
  },
  command: {
    find: 'practicafinal',
    filter: { name: 'Nobuko Linzey' },
    projection: { _id: 0 },
    $sdb: 'lmf'
  }
}

```

```

},
command: {
  find: 'practicafinal',
  filter: { name: 'Nobuko Linzey' },
  projection: { _id: 0 },
  $sdb: 'lmf'
},
serverInfo: {
  host: '13e6561c38af',
  port: 27017,
  version: '6.0.1',
  gitVersion: '32f0f9c88dc44a2c8073a5bd47cf779d4bfbdee6b'
},
serverParameters: {
  internalQueryFacetBufferSizeBytes: 104857600,
  internalQueryFacetMaxOutputDocSizeBytes: 104857600,
  internalLookupStageIntermediateDocumentMaxSizeBytes: 104857600,
  internalDocumentSourceGroupMaxMemoryBytes: 104857600,
  internalQueryMaxBlockingSortMemoryUsageBytes: 104857600,
  internalQueryProhibitBlockingMergeOnMongoS: 0,
  internalQueryMaxAddToSetBytes: 104857600,
  internalDocumentSourceSetWindowFieldsMaxMemoryBytes: 104857600
},
ok: 1
}

```



Lo importante:

nReturned: 1 (devuelve una salida)

TotalKeysExamined: 1 (1 índice)

totalDocsExamined: 1 (examinó 1 documentos y devolvió 1)

Como podemos ver el rendimiento ha mejorado lo máximo, antes revisó 200 documentos para devolver 1 y ahora solo ha tenido que revisar uno y devolver uno. Hemos maximizado el rendimiento con el índice.

## SIN INDICE

- Otra búsqueda diferente

```
inf> db.practicafinal.find({"name":"Cody Strouth"},{"scores":{"$elemMatch":{"type":"quiz"}}},{"scores.score":1})
[
  { _id: 69, scores: [ { score: 99.80348240553108, type: 'quiz' } ] },
  { _id: 183, scores: [ { score: 78.61720316992681, type: 'quiz' } ] }
]
```

Calculamos rendimiento:

```
    },
    rejectedPlans: []
  },
  executionStats: {
    executionSuccess: true,
    nReturned: 2,
    executionTimeMillis: 0,
    totalKeysExamined: 0,
    totalDocsExamined: 200,
    executionStages: {
      stage: 'PROJECTION_DEFAULT',
      nReturned: 2,
      executionTimeMillisEstimate: 0,
      works: 202,
      advanced: 2,
      needTime: 199,
      needYield: 0,
      saveState: 0,
      restoreState: 0,
      isEOF: 1,
      transformBy: { scores: { '$elemMatch': { type: 'quiz' } } },
      inputStage: {
        stage: 'COLLSCAN',
        filter: { name: { '$eq': 'Cody Strouth' } },
        nReturned: 2,
        executionTimeMillisEstimate: 0,
        works: 202,
        advanced: 2,
        needTime: 199,
        needYield: 0,
        saveState: 0,
        restoreState: 0,
        isEOF: 1,
        direction: 'forward',
        docsExamined: 200
      }
    }
  }
}
```



Lo importante aquí:

nReturned: 2 (devuelve dos salidas)

TotalKeysExamined: 0 (no usa índice)

totalDocsExamined: 200 (examinó 200 documentos y devolvió 2)

## CON INDICE

Seguimos con el name como índice

```
imf> db.practicafinal.find({"name":"Cody Strouth"},{"scores":{"$elemMatch":{"type":"quiz"}}},{"scores.score":1}).explain("executionStats")
```

```
imf> db.practicafinal.find({"name":"Cody Strouth"},{"scores":{"$elemMatch":{"type":"quiz"}}},{"scores.score":1}).explain("executionStats")
{
  explainVersion: '1',
  queryPlanner: {
    namespace: 'imf.practicafinal',
    indexFilterSet: false,
    parsedQuery: { name: { '$eq': 'Cody Strouth' } },
    queryHash: '64908032',
    planCacheKey: 'A0C0273F',
    maxIndexedOrSolutionsReached: false,
    maxIndexedAndSolutionsReached: false,
    maxScansToExplodeReached: false,
    winningPlan: {
      stage: 'PROJECTION_DEFAULT',
      transformBy: { scores: { '$elemMatch': { type: 'quiz' } } },
      inputStage: {
        stage: 'FETCH',
        inputStage: {
          stage: 'IXSCAN',
          keyPattern: { name: 1 },
          indexName: 'name_1',
          isMultiKey: false,
          multiKeyPaths: { name: [] },
          isUnique: false,
          isSparse: false,
          isPartial: false,
          indexVersion: 2,
          direction: 'forward',
          indexBounds: { name: [ ['Cody Strouth', 'Cody Strouth'] ] }
        }
      }
    },
    rejectedPlans: []
  },
  executionStats: {
    executionSuccess: true,
    nReturned: 2,
    executionTimeMillis: 0,
    totalKeysExamined: 2,
    totalDocsExamined: 2,
    executionStages: {
      stage: 'PROJECTION_DEFAULT',
      nReturned: 2,
      executionTimeMillisEstimate: 0,
      works: 3,
      advanced: 2,
      needTime: 0,
      needYield: 0
    }
  }
}
```

```

transformBy: { scores: { 'SeleniumMatch': { type: 'quiz' } } },
inputStage: {
  stage: 'FETCH',
  nReturned: 2,
  executionTimeMillisEstimate: 0,
  works: 3,
  advanced: 2,
  needTime: 0,
  needYield: 0,
  saveState: 0,
  restoreState: 0,
  isEOF: 1,
  docsExamined: 2,
  alreadyHasObj: 0,
  inputStage: {
    stage: 'IXSCAN',
    nReturned: 2,
    executionTimeMillisEstimate: 0,
    works: 3,
    advanced: 2,
    needTime: 0,
    needYield: 0,
    saveState: 0,
    restoreState: 0,
    isEOF: 1,
    keyPattern: { name: 1 },
    indexName: 'name_1',
    isMultiKey: false,
    multikeyPaths: { name: [] },
    isUnique: false,
    isSparse: false,
    isPartial: false,
    indexVersion: 2,
    direction: 'forward',
    indexBounds: { name: [ [ 'Cody Strouth', 'Cody Strouth' ] ] },
    keysExamined: 2,
    seeks: 1,
    dupsTested: 0,
    dupsDropped: 0
  }
}
},
command: {
  find: 'practicafinal',
  filter: { name: 'Cody Strouth' },
  projection: { scores: { 'SeleniumMatch': { type: 'quiz' } } },
  '$db': 'inf'
},

```

```

}
},
command: {
  find: 'practicafinal',
  filter: { name: 'Cody Strouth' },
  projection: { scores: { 'SeleniumMatch': { type: 'quiz' } } },
  '$db': 'inf'
},
serverInfo: {
  host: '13e6561c38af',
  port: 27017,
  version: '6.0.1',
  gitVersion: '32f0f9c88dc44a2c8073a5bd47cf779d4bfbdee6b'
},
serverParameters: {
  internalQueryFacetBufferSizeBytes: 104857600,
  internalQueryFacetMaxOutputDocSizeBytes: 104857600,
  internalLookupStageIntermediateDocumentMaxSizeBytes: 104857600,
  internalDocumentSourceGroupMaxMemoryBytes: 104857600,
  internalQueryMaxBlockingSortMemoryUsageBytes: 104857600,
  internalQueryProhibitBlockingMergeOnMongo5: 0,
  internalQueryMaxAddToSetBytes: 104857600,
  internalDocumentSourceSetWindowFieldsMaxMemoryBytes: 104857600
},
ok: 1
}
inf>

```

totalDocsExamined: 2 (examinó 2 documentos)

Como podemos ver el rendimiento ha mejorado lo máximo, antes revisó 200 documentos para devolver 1 y ahora solo ha tenido que revisar uno y devolver uno. Hemos maximizado el rendimiento con el índice.

### 3.2. Operaciones de agregación

## SIN INDICE

```
imf>db.practicafinal.aggregate([{"$unwind": "$scores"}, {"$group": {"_id": "$name", "media": {"$avg": "$scores.score"}}, {"$sort": {"media": -1}}, {"$limit": 1}])
```

```
[ { id: 'Jessika Dagenais', media: 89.89140723964316 } ]
```

### Para ver el rendimiento

```

tnf> db.practicafinal.aggregate([{$unwind: '$scores'},{$group: {'_id': '$name', 'media': {$avg: '$scores.score'}}},{$sort: {'media': 1}},{$limit: 1}]).explain('executionStats')
{
  explainVersion: 1,
  stages: [
    {
      $scan: {
        queryPlanner: {
          namespace: 'not.practicafinal',
          indexFilterSet: false,
          parsedQuery: {},
          queryHash: 'c15f6706',
          planCacheKey: 'c15f6706',
          maxIndexedOrSolutionsReached: false,
          maxIndexedAndSolutionsReached: false,
          maxScansToExplodeReached: false,
          winningPlan: {
            stage: 'PROJECTION_SAMPLE',
            transformBy: { name: 1, scores: 1, _id: 0 },
            inputStage: { stage: 'COLLSCAN', direction: 'forward' }
          },
          rejectedPlans: []
        },
        executionStats: {
          executionSuccess: true,
          nReturned: 280,
          executionTimeMillis: 0,
          totalKeysExamined: 0,
          totalDocsExamined: 280,
          executionStages: {
            stage: 'PROJECTION_SAMPLE',
            nReturned: 280,
            executionTimeMillisEstimate: 0,
            works: 282,
            advanced: 280
          }
        }
      }
    ]
  }
}

```

```
needTime: 199,
needYield: 0,
saveState: 0,
restoreState: 0,
isEOF: 1,
transformBy: { scores: { '$elemMatch': { type: 'quiz' } } },
inputStage: {
  stage: 'COLLSCAN',
  filter: { name: { '$eq': 'Cody Strouth' } },
  nReturned: 2,
  executionTimeMillisEstimate: 0,
  works: 202,
  advanced: 2,
  needTime: 199,
  needYield: 0,
```

Lo importante aquí:

nReturned: 200 (devuelve)

TotalKeysExamined: 0 (no usa índice)

totalDocsExamined: 200 (examinó 200 documentos)

## CON INDICE

Uso como índice scores

```
imf> db.inventory.createIndex( { scores: 1 } )
```

scores\_1

```
imf> db.practicafinal.aggregate([{"$unwind":"$scores"},{$group: { _id: "$name",  
"media": { $avg: "$scores.score" } }},{$sort: { media:  
-1 }},{$limit:1}]).explain("executionStats")
```

```
namespace: "imf.practicafinal",  
indexFilterSet: false,  
parsedQuery: {},  
queryHash: 'C15F67D0',  
planCacheKey: 'C15F67D0',  
maxIndexedOrSolutionsReached: false,  
maxIndexedAndSolutionsReached: false,  
maxScansToExplodeReached: false,  
winningPlan: {  
  stage: 'PROJECTION_SIMPLE',  
  transformBy: { name: 1, scores: 1, _id: 0 },  
  inputStage: { stage: 'COLLSCAN', direction: 'forward' }  
},  
rejectedPlans: []  
},  
executionStats: {  
  executionSuccess: true,  
  nReturned: 200,  
  executionTimeMillis: 1,  
  totalKeysExamined: 0,  
  totalDocsExamined: 200,  
  executionStages: {  
    stage: 'PROJECTION_SIMPLE',  
    nReturned: 200,  
    executionTimeMillisEstimate: 0,  
    works: 202,  
    advanced: 200,  
    needTime: 1,  
    needYield: 0,  
    saveState: 1,  
    restoreState: 1,  
    isEOF: 1,  
    transformBy: { name: 1, scores: 1, _id: 0 },  
    inputStage: {  
      stage: 'COLLSCAN',  
      nReturned: 200,  
      executionTimeMillisEstimate: 0,  
      works: 202,  
      advanced: 200,  
      needTime: 1,  
      needYield: 0,  
      saveState: 1,  
      restoreState: 1,  
      isEOF: 1,  
      direction: 'forward',  
      docsExamined: 200  
    }  
  }  
}
```



```

{
  'Sunwind': { path: '$scores' },
  nReturned: Long("600"),
  executionTimeMillisEstimate: Long("0")
},
{
  '$group': { _id: '$name', media: { '$avg': '$scores.score' } },
  maxAccumulatorMemoryUsageBytes: { media: Long("10032") },
  totalOutputDataSizeBytes: Long("28180"),
  usedDisk: false,
  spills: Long("0"),
  nReturned: Long("114"),
  executionTimeMillisEstimate: Long("0")
},
{
  '$sort': { sortKey: { media: -1 }, limit: Long("1") },
  totalDataSizeSortedBytesEstimate: Long("0"),
  usedDisk: false,
  spills: Long("0"),
  nReturned: Long("1"),
  executionTimeMillisEstimate: Long("0")
}
],
serverInfo: {
  host: '13e6561c38af',
  port: 27017,
  version: '6.0.1',
  gitVersion: '32f0f9c88dc44a2c8073a5bd47cf779d4b7dee6b'
},
serverParameters: {
  internalQueryFacetBufferSizeBytes: 104857600,
  internalQueryFacetMaxOutputDocSizeBytes: 104857600,
  internalLookupStageIntermediateDocumentMaxSizeBytes: 104857600,
  internalDocumentSourceGroupMaxMemoryBytes: 104857600,
  internalQueryMaxBlockingSortMemoryUsageBytes: 104857600,
  internalQueryProhibitBlockingMergeOnMongoS: 0,
  internalQueryMaxAddToSetBytes: 104857600,
  internalDocumentSourceSetWindowFieldsMaxMemoryBytes: 104857600
},
command: {
  aggregate: 'practicafinal',
  pipeline: [
    { 'Sunwind': '$scores' },
    {
      '$group': { _id: '$name', media: { '$avg': '$scores.score' } }
    },
    { '$sort': { media: -1 } },
    { '$limit': 1 }
  ]
}

```

Lo importante aquí:

nReturned: 200 (devuelve)

TotalKeysExamined: 0 (no usa el índice)

totalDocsExamined: 200 (examinó 200 documentos)

Como podemos ver el rendimiento es igual en aggregate, con o sin índice

Y si sólo uno como índice name se llega al mismo resultado

```
    stage: 'PROJECTION_SIMPLE',
    transformBy: { name: 1, scores: 1, _id: 0 },
    inputStage: { stage: 'COLLSCAN', direction: 'forward' }
  },
  rejectedPlans: []
},
executionStats: {
  executionSuccess: true,
  nReturned: 200,
  executionTimeMillis: 1,
  totalKeysExamined: 0,
  totalDocsExamined: 200,
  executionStages: {
    stage: 'PROJECTION_SIMPLE',
    nReturned: 200,
    executionTimeMillisEstimate: 0,
    works: 202,
    advanced: 200,
    needTime: 1,
    needYield: 0,
    saveState: 1,
    restoreState: 1,
    isEOF: 1,
    transformBy: { name: 1, scores: 1, _id: 0 },
    inputStage: {
      stage: 'COLLSCAN',
      nReturned: 200,
      executionTimeMillisEstimate: 0,
      works: 202,
      advanced: 200,
      needTime: 1,
      needYield: 0,
      saveState: 1,
      restoreState: 1,
      isEOF: 1,
      direction: 'forward',
      docsExamined: 200
    }
  }
}
```

Y si usamos los dos índices a la vez, no reconoce la función con aggregate:



```
imf> db.practicafinal.aggregate([{"$unwind":"$scores"},{$group: {_id: "$name",
"media": {$avg: "$scores.score"}}},{$sort: {media: -1}},{$limit:1}]).hint({ name: 1,
scores: 1 }).explain("executionStats")
```

**TypeError: db.practicafinal.aggregate(...).hint is not a function**

Como podemos ver el rendimiento es igual en aggregate, con o sin índice

Vemos otro ejemplo de agregación:

### SIN INDICE

```
imf> db.practicafinal.aggregate([{"$unwind":"$scores"},{$group:{_id:"$scores.type",
"media":{$avg: "$scores.score"}}}])
```

```
[
  { _id: 'exam', media: 48.67367075950175 },
  { _id: 'quiz', media: 48.99672319430254 },
  { _id: 'homework', media: 67.81869620661149 }
```

```
imf> db.practicafinal.aggregate([{"$unwind":"$scores"},{$group: {_id: "$scores.type", "media": {$avg: "$scores.score"}}}]).explain("executionStats")
{
  explainVersion: '1',
  stages: [
    {
      'cursor': {
        queryPlanner: {
          namespace: 'imf.practicafinal',
          indexFilterSet: false,
          parsedQuery: {},
          queryHash: '528CE6EA',
          planCacheKey: '528CE6EA',
          maxIndexedOrSolutionsReached: false,
          maxIndexedAndSolutionsReached: false,
          maxScansToExplodeReached: false,
          winningPlan: {
            stage: 'PROJECTION_SIMPLE',
            transformBy: [ scores: 1, _id: 0 ],
            inputStage: { stage: 'COLLSCAN', direction: 'forward' }
          },
          rejectedPlans: []
        },
        executionStats: {
          executionSuccess: true,
          nReturned: 200,
          executionTimeMillis: 1,
          totalKeysExamined: 0,
          totalDocsExamined: 200,
          executionStages: {
            stage: 'PROJECTION_SIMPLE',
            nReturned: 200,
            executionTimeMillisEstimate: 0,
            works: 202,
            advanced: 200,
            needTime: 1,
            needYield: 0,
            saveState: 1,
            restoreState: 1,
            isEOF: 1,
            transformBy: [ scores: 1, _id: 0 ],
            inputStage: {
              stage: 'COLLSCAN'
```

Lo importante aquí:

nReturned: 200 (devuelve)

TotalKeysExamined: 0 (sin índice)

totalDocsExamined: 200 (examinó 200 documentos)

## CON INDICE

```
imf> db.practicafinal.aggregate([{"$unwind":"$scores"},{$group:{_id:"$scores.type",
"media":{$avg: "$scores.score"}}}])
```

```
    transformBy: { scores: 1, _id: 0 },
    inputStage: { stage: 'COLLSCAN', direction: 'forward' }
  },
  rejectedPlans: []
},
executionStats: {
  executionSuccess: true,
  nReturned: 200,
  executionTimeMillis: 1,
  totalKeysExamined: 0,
  totalDocsExamined: 200,
  executionStages: {
    stage: 'PROJECTION_SIMPLE',
    nReturned: 200,
    executionTimeMillisEstimate: 0,
    works: 202,
    advanced: 200,
    needTime: 1,
    needYield: 0,
    saveState: 1,
    restoreState: 1,
    isEOF: 1,
    transformBy: { scores: 1, _id: 0 },
    inputStage: {
      stage: 'COLLSCAN',
      nReturned: 200,
      executionTimeMillisEstimate: 0,
      works: 202,
      advanced: 200,
      needTime: 1,
      needYield: 0,
      saveState: 1,
      restoreState: 1,
      isEOF: 1,
      direction: 'forward',
      docsExamined: 200
    }
  }
},
nReturned: Long("200"),
executionTimeMillisEstimate: Long("0")
},
{
  '$unwind': { path: '$scores' },
  nReturned: Long("600"),
  executionTimeMillisEstimate: Long("0")
},
{
```

Como podemos ver el rendimiento es igual en aggregate, con o sin índice

### 3.3. Operaciones de actualización

#### CON/SIN INDICE

```
imf>db.practicafinal.updateOne({"name":"Nobuko Linzey"},{ $set: {
"scores.$[elem].score" : 80.5 }},{ arrayFilters: [ { "elem.type": "exam" } ] })
```

```
imf> db.practicafinal.find({"name":"Nobuko Linzey"})
```

```
[ { _id: 110,
  name: 'Nobuko Linzey',
  scores: [
    { score: 21.83, type: 'exam' },
    { score: 58.58331128403415, type: 'quiz' },
    { score: 47.44831568815929, type: 'homework' }
  ]
}]
```

```
imf> db.practicafinal.updateOne({"name":"Nobuko Linzey"},{ $set: {
"scores.$[elem].score" : 21.83 } },{ arrayFilters: [ { "elem.type": "exam" } ]
}).explain("executionStats")
```

**TypeError: db.practicafi ... } ] }).explain is not a function**

No se aplica la función explain en update. No podemos ver el rendimiento así.

### 3. 4. Operaciones de Borrado

#### CON/SIN INDICE

```
db.practicafinal.deleteOne({'name':'Efrain Claw'})
```

```
> db.practicafinal.deleteOne({'name':'Efrain Claw'})  
{ "acknowledged" : true, "deletedCount" : 1 }  
db.practicafinal.deleteOne({'name':'Efrain Claw'}).pretty()
```

```
imf> db.practicafinal.deleteOne({'name':'Efrain Claw'})
```

```
{ acknowledged: true, deletedCount: 1 }
```

```
imf> db.practicafinal.deleteOne({'name':'Efrain Claw'}).explain("executionStats")
```

**TypeError: db.practicafi ... law'}).explain is not a function**

No podemos aplicar explain tampoco en funciones de borrado.